


# ELEC 576: Assignment 1

Wilfredo J. Molina

October 3, 2018



# Contents

<b>1</b>	<b>Big Picture</b>	<b>3</b>
<b>2</b>	<b>Neural Network of 3 Layers</b>	<b>4</b>
2.1	Back-Propagation . . . . .	5
2.2	Experiments . . . . .	6
2.2.1	$n_2 = 10$ , make_moons Data Set . . . . .	7
2.2.2	$n_2 = 20$ , make_moons Data Set . . . . .	8
2.2.3	$n_2 = 30$ , make_moons Data Set . . . . .	9
2.2.4	$n_2 = 10$ , make_circles Data Set . . . . .	10
2.2.5	$n_2 = 20$ , make_circles Data Set . . . . .	11
2.2.6	$n_2 = 30$ , make_circles Data Set . . . . .	12
<b>3</b>	<b>Neural Network of <math>n</math> Layers</b>	<b>13</b>
3.1	Back-Propagation . . . . .	14
3.2	Experiments . . . . .	15
3.2.1	$(n_2, n_3, n_4) = (12, 12, 12)$ , make_moons Data Set . . . . .	16
3.2.2	$(n_2, n_3, n_4) = (16, 8, 16)$ , make_moons Data Set . . . . .	17
3.2.3	$(n_2, n_3, n_4) = (10, 16, 10)$ , make_moons Data Set . . . . .	18
3.2.4	$(n_2, n_3, n_4) = (12, 12, 12)$ , make_circles Data Set . . . . .	19
3.2.5	$(n_2, n_3, n_4) = (16, 8, 16)$ , make_circles Data Set . . . . .	20
3.2.6	$(n_2, n_3, n_4) = (10, 16, 10)$ , make_circles Data Set . . . . .	21
<b>4</b>	<b>Convolutional Neural Network</b>	<b>22</b>
<b>5</b>	<b>Definitions and Remarks</b>	<b>29</b>

# 1 Big Picture

Suppose that we have a set  $\{x_i\}_{i=1}^N \subset \mathbb{R}^{n_1}$  of  $N \in \mathbb{N}$  *data points*, each with  $n_1 \in \mathbb{N}$  *dimensions*, and that we have *classified* each such point as belonging to one of  $n_L \in \mathbb{N}$  *classes*.

Let  $\{y_i\}_{i=1}^N \subset \mathcal{C}^{n_L}$  and note that the information that  $x_i$  belongs to a particular class is captured by the ordered pair  $(x_i, y_i)$ .

Our goal is to find a function

$$\begin{aligned} \Psi : \mathbb{R}^{n_1} \rightarrow (0, 1)^{n_L} \text{ such that } -\frac{1}{N} \sum_{i=1}^N \langle y_i, \ln(\Psi(x_i)) \rangle \text{ is } \textit{close enough} \text{ to } 0 \text{ and that} \\ \langle \Psi(x), \underbrace{(1, 1, \dots, 1)}_{n_L \text{ times}} \rangle = 1 \text{ for every } x \in \mathbb{R}^{n_1}. \end{aligned} \quad (1.1)$$

In doing so, we hope that  $\Psi$  can help us classify points outside of  $\{x_i\}_{i=1}^N$  in a manner that is consistent with our expectations, vague as they may be.

Note that  $\mathcal{C}^{n_L}$  and the co-domain of  $\Psi$  are chosen in this way to represent probabilities and that we implicitly extended  $\ln : (0, \infty) \rightarrow \mathbb{R}$  above to  $\ln : (0, \infty)^{n_L} \rightarrow \mathbb{R}^{n_L}$  point-wise, i.e.,  $\ln(x_1, x_2, \dots, x_{n_L}) = (\ln(x_1), \ln(x_2), \dots, \ln(x_{n_L}))$ . We will continue to do this as it should be unambiguous in this context.

## 2 Neural Network of 3 Layers

$$\begin{aligned} \Psi : \mathbb{R}^{n_1} &\rightarrow (0, 1)^{n_3}, \text{ defined by} \\ x &\mapsto \sigma_{n_3} (W_2 (\varphi (W_1 (x) + b_1)) + b_2), \end{aligned} \tag{2.1}$$

is a *neural network of 3 layers*, where  $W_i \in \text{Hom}(\mathbb{R}^{n_i}, \mathbb{R}^{n_{i+1}})$ ,  $b_i \in \mathbb{R}^{n_{i+1}}$ ,  $\varphi \in \mathcal{D}$ , and  $n_i \in \mathbb{N}$ .

Given data points  $\{x_i\}_{i=1}^N \subset \mathbb{R}^{n_1}$  and classes  $\{y_i\}_{i=1}^N \subset \mathcal{C}^{n_3}$ , we expect  $\Psi$  to be a *good enough* candidate for (1.1) for a suitable choice of  $W_i$ ,  $b_i$ , and  $\varphi$ .

We justify our requirement that  $\varphi$  be in  $\mathcal{D}$  by stating that it is *unlikely* that we will evaluate it at the points where it is not differentiable if we initialize  $W_i$  and  $b_i$  *well enough*. However, note that  $\varphi : \mathbb{R} \rightarrow \mathbb{R}$ , defined by  $x \mapsto \sum_{n \in \mathbb{N}} |\sin(n\pi x)| / n^3$ , is differentiable only on  $\mathbb{R} \setminus \mathbb{Q}$ , i.e.,  $\varphi \in \mathcal{D}$ , but no machine can represent the points where  $\varphi$  is differentiable. Therefore, in practice, our requirement translates to  $\varphi$  being a function such that  $|D \cap R|$  is *small enough*, where  $D$  is the set of points where  $\varphi$  is not differentiable and  $R \subset \mathbb{R}$  is the set of numbers that a particular machine can represent (after all,  $R$  has Lebesgue measure 0 for every machine that has ever been and will be made).

## 2.1 Back-Propagation

Given data points  $\{x_i\}_{i=1}^N \subset \mathbb{R}^{n_1}$  and classes  $\{y_i\}_{i=1}^N \subset \mathcal{C}^{n_3}$ , we fix  $\varphi \in \mathcal{D}$ , initialize  $W_i$  and  $b_i$  *well enough*, and attempt to find a suitable, *a posteriori* choice of  $W_i$  and  $b_i$  for (2.1) via an algorithm called *back-propagation*. To this end, we define the *loss function*

$$E : \mathbb{R}^{n_1} \times \mathcal{C}^{n_3} \times \text{Hom}(\mathbb{R}^{n_1}, \mathbb{R}^{n_2}) \times \mathbb{R}^{n_2} \times \text{Hom}(\mathbb{R}^{n_2}, \mathbb{R}^{n_3}) \times \mathbb{R}^{n_3} \rightarrow \mathbb{R} \text{ by} \quad (2.2)$$

$$(x, y, W_1, b_1, W_2, b_2) \mapsto -\langle y, \ln(\sigma_{n_3}(W_2(\varphi(W_1(x) + b_1)) + b_2)) \rangle$$

and, in the interest of implementation, fix  $j \in \{1, 2, \dots, N\}$  and let

$$\begin{aligned} a_0 &:= x_j, \\ z_1 &:= W_1(a_0) + b_1, \\ a_1 &:= \varphi(z_1), \\ z_2 &:= W_2(a_1) + b_2, \\ a_2 &:= \sigma_{n_3}(z_2), \text{ and} \\ \delta_i &:= \nabla_{z_i} E. \end{aligned}$$

Then

$$\begin{aligned} \delta_2 &:= a_2 - y \text{ and} \\ \delta_1 &:= W_2^*(\delta_2) \odot \varphi'(z_1). \end{aligned}$$

Letting  $\eta > 0$  be *close enough* to 0 (we call  $\eta$  the *learning rate*), we hope that repeating the assignments

$$\begin{aligned} W_i &\leftarrow W_i - \eta \overline{a_{i-1}} \otimes \delta_i \text{ and} \\ b_i &\leftarrow b_i - \eta \delta_i \end{aligned}$$

*enough times* for every  $j \in \{1, 2, \dots, N\}$  will *eventually* yield a *good enough* candidate for (1.1), which we consider to be *trained*.

## 2.2 Experiments

Each picture in the next few pages is a visual representation of 1000 distinct elements of some square  $[a, b] \times [c, d] \subset \mathbb{R}^2$ , each belonging to one of two classes (red or blue), where 0.000, 0.025, and 0.075 represent different levels of *noise* (0.000 represents no noise), and of the *decision boundary* of a trained neural network of 3 layers with  $n_1 = 2$ ,  $n_2 \in \{10, 20, 30\}$ ,  $n_3 = 2$ ,

$$\varphi(x) = \begin{cases} 0 & \text{if } x < 0 \text{ and} \\ x & \text{if } x \geq 0, \end{cases} \quad (\text{ReLU})$$


$$\varphi(x) = \frac{1}{1 + e^{-x}}, \text{ and} \quad (\text{sigmoid})$$

$$\varphi(x) = \tanh(x). \quad (\text{hyperbolic tangent})$$


We implement  $\Psi$  in Python as a class that inherits all of its fields and methods from the class that we implemented for (3.1): the following is the entire code for this class.

```
class ThreeLayerNeuralNetwork(DeepNeuralNetwork):
    def __init__(self,
                  input_dim = 2,
                  hidden_dim = 10,
                  output_dim = 2,
                  activation_type = 'relu',
                  regularization = 0,
                  random_seed = None):
        super().__init__([input_dim,
                           hidden_dim,
                           output_dim],
                           activation_type,
                           regularization,
                           random_seed)
```

### 2.2.1 $n_2 = 10$ , make\_moons Data Set




### 2.2.2 $n_2 = 20$ , make\_moons Data Set







### 2.2.3 $n_2 = 30$ , make\_moons Data Set




### 2.2.4 $n_2 = 10$ , make\_circles Data Set



### 2.2.5 $n_2 = 20$ , make\_circles Data Set



### 2.2.6 $n_2 = 30$ , make\_circles Data Set



### 3 Neural Network of $n$ Layers

$$\begin{aligned} \Psi : \mathbb{R}^{n_1} &\rightarrow (0, 1)^{n_L}, \text{ defined by} \\ x &\mapsto \sigma_{n_L} (W_{L-1} (\varphi (W_{L-2} (\cdots \varphi (W_1 (x) + b_1) \cdots) + b_{L-2})) + b_{L-1}), \end{aligned} \quad (3.1)$$

is a *neural network* of  $L \in \mathbb{N}$  layers, where  $W_i \in \text{Hom}(\mathbb{R}^{n_i}, \mathbb{R}^{n_{i+1}})$ ,  $b_i \in \mathbb{R}^{n_{i+1}}$ ,  $\varphi \in \mathcal{D}$ ,  $n_i \in \mathbb{N}$ , and  $L > 1$ .

Given data points  $\{x_i\}_{i=1}^N \subset \mathbb{R}^{n_1}$  and classes  $\{y_i\}_{i=1}^N \subset \mathcal{C}^{n_L}$ , we expect  $\Psi$  to be a *better* candidate than (2.1) for (1.1) for a suitable choice of  $W_i$ ,  $b_i$ ,  $L$ , and  $\varphi$ .

### 3.1 Back-Propagation

Given data points  $\{x_i\}_{i=1}^N \subset \mathbb{R}^{n_1}$  and classes  $\{y_i\}_{i=1}^N \subset \mathcal{C}^{n_3}$ , we fix  $\varphi \in \mathcal{D}$ , initialize  $W_i$  and  $b_i$  *well enough*, and attempt to find a suitable, *a posteriori* choice of  $W_i$  and  $b_i$  for (3.1). To this end, we fix  $j \in \{1, 2, \dots, N\}$  and let

$$\begin{aligned} a_0 &:= x_j, \\ z_i &:= W_i(a_{i-1}) + b_i, \\ a_i &:= \varphi(z_i), \\ a_{L-1} &:= \sigma_{n_L}(z_{L-1}), \text{ and} \\ \delta_i &:= \nabla_{z_i} E, \end{aligned}$$

where the definition of  $E$  is analogous to that of (2.2). Then

$$\begin{aligned} \delta_{L-1} &:= a_{L-1} - y \text{ and} \\ \delta_i &:= W_{i+1}^*(\delta_{i+1}) \odot \varphi'(z_i). \end{aligned}$$

Letting  $\eta > 0$  be *close enough* to 0, we hope that repeating the assignments

$$\begin{aligned} W_i &\leftarrow W_i - \eta \overline{a_{i-1}} \otimes \delta_i \text{ and} \\ b_i &\leftarrow b_i - \eta \delta_i, \end{aligned}$$

*enough times* for every  $j \in \{1, 2, \dots, N\}$  will *eventually* yield a *good enough* candidate for (1.1).

## 3.2 Experiments

Each picture in the next few pages is a visual representation of 1000 distinct elements of some square  $[a, b] \times [c, d] \subset \mathbb{R}^2$ , each belonging to one of two classes (red or blue), where 0.000, 0.025, and 0.075 represent different levels of *noise* (0.000 represents no noise), and of the *decision boundary* of a trained neural network of 5 layers with  $n_1 = 2$ ,  $(n_2, n_3, n_4) \in \{(12, 12, 12), (16, 8, 16), (10, 16, 10)\}$ ,  $n_5 = 2$ ,

$$\varphi(x) = \begin{cases} 0 & \text{if } x < 0 \text{ and} \\ x & \text{if } x \geq 0, \end{cases} \quad (\text{ReLU})$$

$$\varphi(x) = \frac{1}{1 + e^{-x}}, \text{ and} \quad (\text{sigmoid})$$

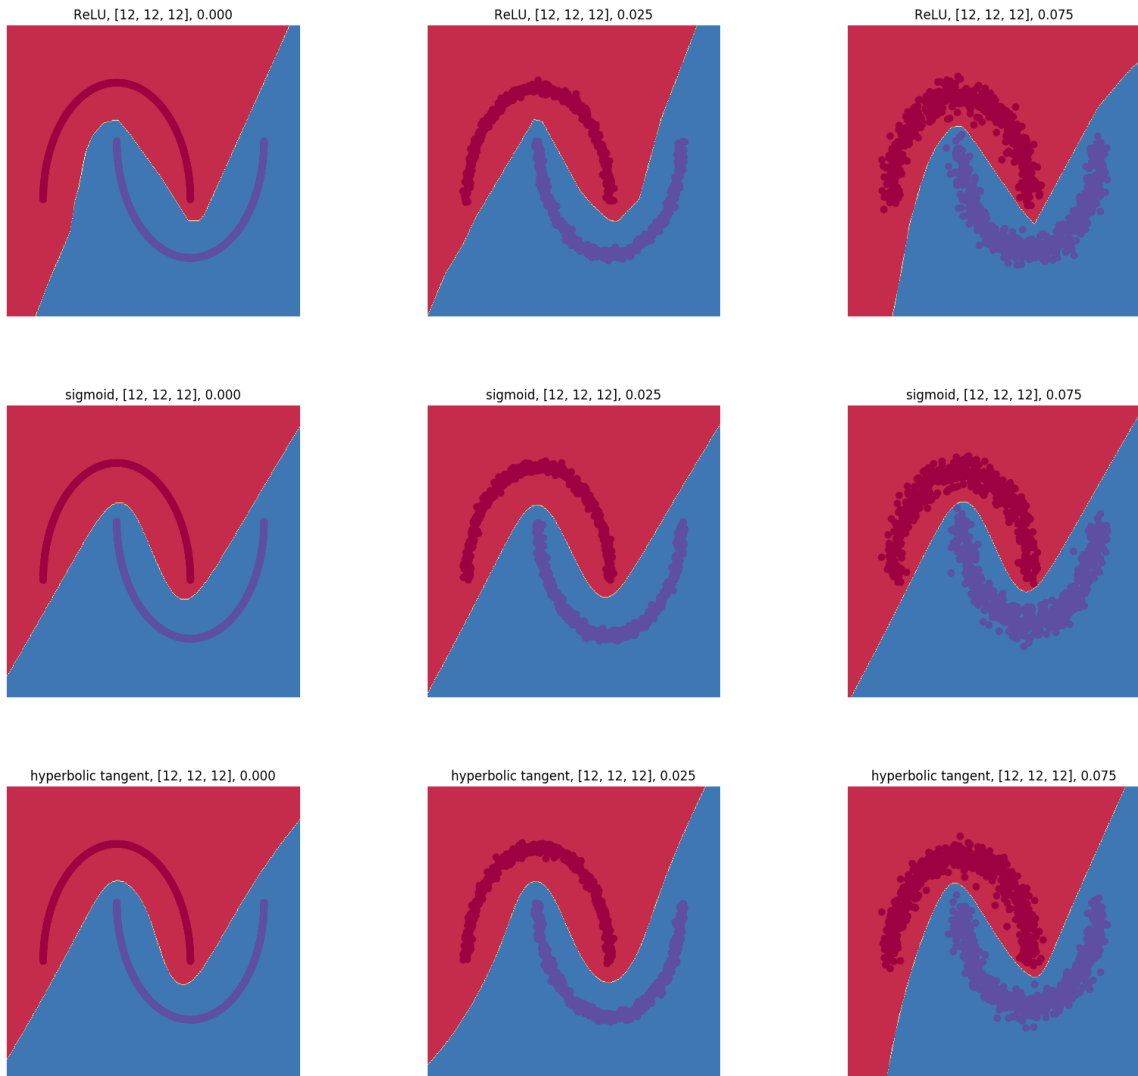
$$\varphi(x) = \tanh(x). \quad (\text{hyperbolic tangent})$$

We implement  $\Psi$  in Python as a class with the following fields and methods:

```
self.dimensions      # NumPy Array
self.activation_type  # String
self.regularization  # Number
self.W               # Dictionary of NumPy Arrays
self.b               # Dictionary of NumPy Arrays
self.a               # Dictionary of NumPy Arrays
self.z               # Dictionary of NumPy Arrays


def __init__(self,
              dimensions = np.array([2, 10, 2]),
              activation_type = 'relu',
              regularization = 0,
              random_seed = None)
def feed_forward(self, a1)
def activation(self, x)
def activation_derivative(self, x)
def soft_max(x)
def back_propagation(self, a1, y)
def train(self, a1, y, train_rate, passes, print_loss, print_rate)
def calculate_loss(self, a1, y)
def predict(self, a1)
```

### 3.2.1 $(n_2, n_3, n_4) = (12, 12, 12)$ , make\_moons Data Set







### 3.2.2 $(n_2, n_3, n_4) = (16, 8, 16)$ , make\_moons Data Set




### 3.2.3 $(n_2, n_3, n_4) = (10, 16, 10)$ , **make\_moons** Data Set




### 3.2.4 $(n_2, n_3, n_4) = (12, 12, 12)$ , **make\_circles** Data Set



### 3.2.5 $(n_2, n_3, n_4) = (16, 8, 16)$ , make\_circles Data Set



### 3.2.6 $(n_2, n_3, n_4) = (10, 16, 10)$ , **make\_circles** Data Set




## 4 Convolutional Neural Network

We implement a convolutional neural network with the following architecture:

```
Convolution(5-5-1-32) - ReLU - MaxPool(2-2) - Convolution(5-5-1-64) - ReLU  
- MaxPool(2-2) - Flatten(1024) - ReLU - Dropout(.5) - SoftMax(10)
```

and train it on MNIST, which is a database of 60000  $28 \times 28 \times 1$  images that look like this:



Using Keras, this can be done as follows:

```
model = Sequential()  
model.add(Conv2D(32, kernel_size = (5, 5), activation = 'relu', input_shape = input_shape))  
model.add(MaxPooling2D((2, 2)))  
model.add(Conv2D(64, (5, 5), activation = 'relu'))  
model.add(MaxPooling2D((2, 2)))  
model.add(Flatten())  
model.add(Dense(1024, activation = 'relu'))  
model.add(Dropout(0.5))  
model.add(Dense(10, activation = 'softmax'))
```

We train our convolutional neural network for 5 epochs with a batch size of 50. This can be done as follows:

```
(x_train, y_train), (x_test, y_test) = mnist.load_data()


model.compile(loss      = keras.losses.categorical_crossentropy,
              optimizer = keras.optimizers.Adadelta(),
              metrics   = ['accuracy'])

model.fit(x_train,
          y_train,
          batch_size    = 50,
          epochs         = 5,
          validation_data = (x_test, y_test))
```


With this setup, we obtain a validation accuracy of 0.9882.

In the next few pages, we visualize our training using tensorboard.


Layer1/biases/stddev/max\_1




Layer1/biases/stddev/min\_1




Layer1/biases/stddev/stddev




Layer1/biases/summaries/mean\_1




Layer1/weights/stddev/max\_1




Layer1/weights/stddev/min\_1




Layer1/weights/stddev/stddev




Layer1/weights/summaries/mean\_1




Layer2/biases/stddev/max\_1




Layer2/biases/stddev/min\_1




Layer2/biases/stddev/stddev




Layer2/biases/summaries/mean\_1




Layer2/weights/stddev/max\_1



Layer2/weights/stddev/min\_1




Layer2/weights/stddev/stddev







Layer2/weights/summaries/mean\_1




accuracy/accuracy\_1




crossentropy/crossentropy




dropout/dropout\_keep\_probability




flatlayer/biases/stddev/max\_1




flatlayer/biases/stddev/min\_1




Layer2/weights/stddev/max\_1




flatlayer/biases/stddev/stddev




flatlayer/biases/summaries/mean\_1




flatlayer/weights/stddev/max\_1




flatlayer/weights/stddev/min\_1




flatlayer/weights/stddev/stddev




flatlayer/weights/summaries/mean\_1




outputlayer/biases/stddev/max\_1




outputlayer/biases/stddev/min\_1




outputlayer/biases/stddev/stddev




outputlayer/biases/summaries/mean\_1




outputlayer/weights/stddev/max\_1




outputlayer/weights/stddev/min\_1




outputlayer/weights/stddev/stddev




outputlayer/weights/summaries/mean\_1




Layer1/activations




Layer1/biases/stddev/histogram




Layer1/linearity/pre-activations




Layer1/weights/stddev/histogram




Layer2/activations




Layer2/biases/stddev/histogram




Layer2/linearity/pre-activations




Layer2/weights/stddev/histogram




Layer2/weights/stddev/histogram




flatlayer/Wx\_plus\_b/pre\_activations




flatlayer/activations




flatlayer/biases/stddev/histogram




flatlayer/weights/stddev/histogram




outputlayer/Wx\_plus\_b/pre\_activations




outputlayer/biases/stddev/histogram




outputlayer/weights/stddev/histogram




Layer1/activations




Layer1/biases/stddev/histogram




Layer1/linearity/pre-activations




Layer1/weights/stddev/histogram




Layer2/activations




Layer2/biases/stddev/histogram




Layer2/linearity/pre-activations




Layer2/weights/stddev/histogram




flatlayer/Wx\_plus\_b/pre\_activations




flatlayer/activations




flatlayer/biases/stddev/histogram




flatlayer/weights/stddev/histogram




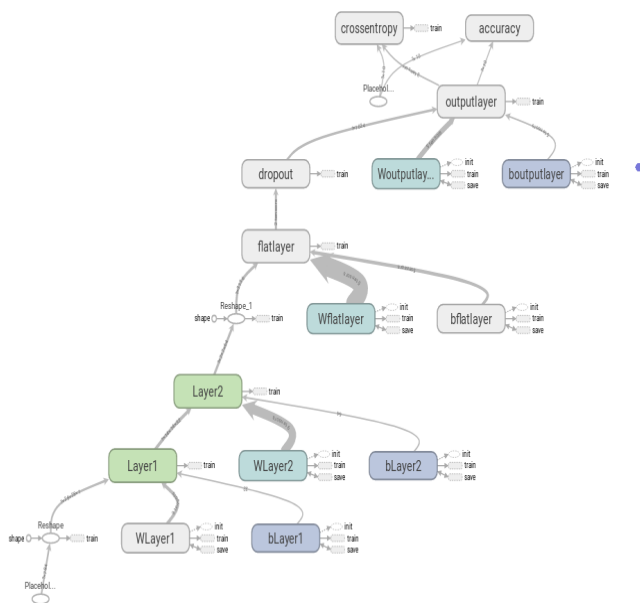

outputlayer/Wx\_plus\_b/pre\_activations



outputlayer/biases/stddev/histogram



outputlayer/weights/stddev/histogram



## 5 Definitions and Remarks

**Definition 5.1.**  $\mathbb{N} := \{1, 2, \dots\}$  is the set of *natural numbers*.

**Definition 5.2.**  $\mathbb{R}$  is the set of *real numbers*.

**Definition 5.3.**  $\mathcal{C}^n := \{(c_1, c_2, \dots, c_n) \in \mathbb{R}^n : c_i = 1 \text{ and } c_j = 0 \text{ if } j \neq i, \text{ where } i \in \{1, 2, \dots, n\}\}$  is the set of *n classes*.

**Definition 5.4.**  $\text{Hom}(V, W)$  is the set of *linear maps*  $V \rightarrow W$ , i.e.,  $T \in \text{Hom}(V, W)$  implies  $T(\lambda v + w) = \lambda T(v) + T(w)$  for every  $v, w \in V$  and every  $\lambda \in \mathcal{F}$ , where  $V$  and  $W$  are linear spaces over a field  $\mathcal{F}$ .

**Remark 5.5.** If  $V$  and  $W$  are finite-dimensional linear spaces (like  $\mathbb{R}^n$ ), then  $T \in \text{Hom}(V, W)$  can be represented by a *matrix*.

**Definition 5.6.**  $V^* := \text{Hom}(V, \mathcal{F})$  is the *dual* of the linear space  $V$  over the field  $\mathcal{F}$ .

**Definition 5.7.**  $T^* : W^* \rightarrow V^*$ , defined by  $f \mapsto f \circ T$ , is the *adjoint* of  $T \in \text{Hom}(V, W)$ , where  $V$  and  $W$  are linear spaces.

**Remark 5.8.** If  $V$  and  $W$  are finite-dimensional linear spaces and  $T \in \text{Hom}(V, W)$ , then  $T^*$  is the *transpose* of  $T$ , when  $T$  is represented by a matrix, i.e.,  $T^* = T^T$ .

**Definition 5.9.**  $v \odot w := (v_1 w_1, v_2 w_2, \dots, v_n w_n)$ , where  $v = (v_1, v_2, \dots, v_n)$ ,  $w = (w_1, w_2, \dots, w_n) \in \mathbb{R}^n$ , is a *Hadamard product*.

**Definition 5.10.**  $\mathcal{D}$  is the set of functions  $\mathbb{R} \rightarrow \mathbb{R}$  that are *differentiable almost everywhere* with respect to the Lebesgue measure on  $\mathbb{R}$ .

**Definition 5.11.**  $f \otimes w : V \rightarrow W$ , defined by  $v \mapsto f(v)w$ , is the *outer product* of  $f$  and  $w$ , where  $f \in V^*$ ,  $w \in W$ , and  $V$  and  $W$  are linear spaces.

**Definition 5.12.**  $\langle \cdot, \cdot \rangle : \mathbb{R}^n \times \mathbb{R}^n \rightarrow \mathbb{R}$ , defined by  $((v_1, v_2, \dots, v_n), (w_1, w_2, \dots, w_n)) \mapsto \sum_{i=1}^n v_i w_i$ , is the *dot product* on  $\mathbb{R}^n$ .

**Definition 5.13.**  $\sigma_n : \mathbb{R}^n \rightarrow \mathbb{R}^n$ , defined by  $(x_1, x_2, \dots, x_n) \mapsto (e^{x_1}, e^{x_2}, \dots, e^{x_n}) / \sum_{i=1}^n e^{x_i}$ , is the *softmax* function on  $\mathbb{R}^n$ .

**Remark/Definition 5.14.** If  $V$  is a finite-dimensional linear space over a field  $\mathcal{F}$ , then  $V$  and  $V^*$  are isomorphic, i.e., there is a bijection  $\phi : V \rightarrow V^*$  such that  $\phi(\lambda v + w) = \lambda \phi(v) + \phi(w)$  for every  $v, w \in V$  and every  $\lambda \in \mathcal{F}$ . In this case, let  $\bar{v} := \phi(v)$ .